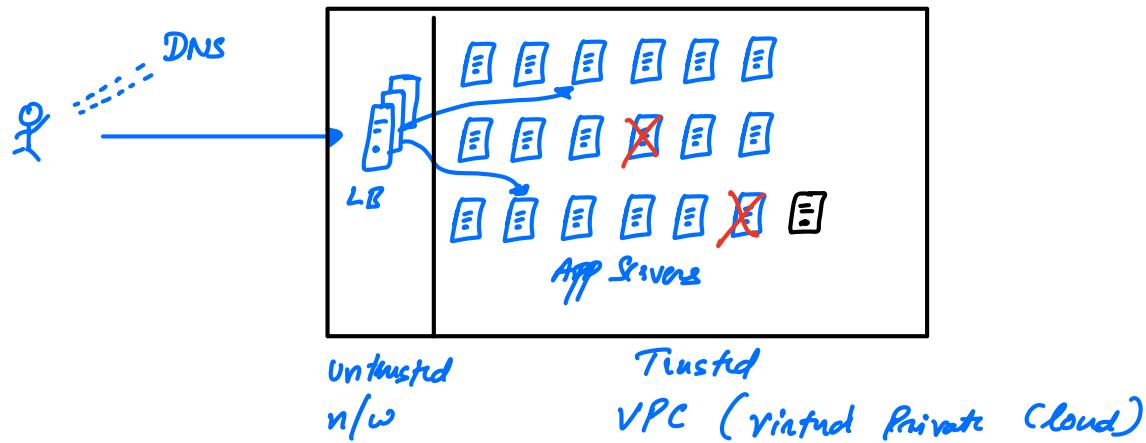


Recap

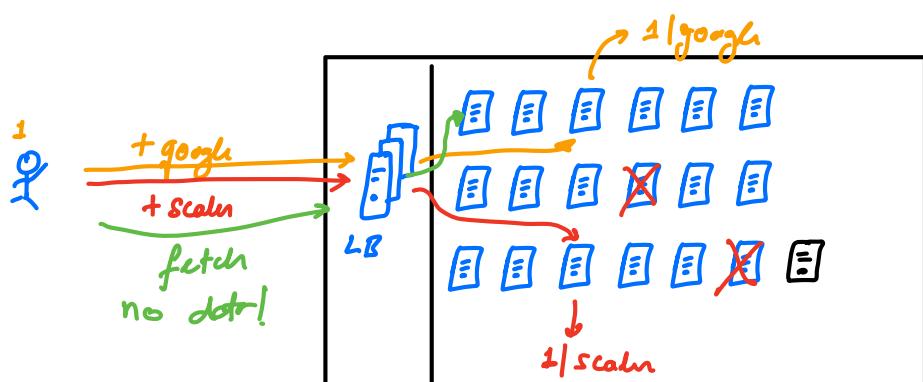
- Case Study - delicious - bookmarking website
- Minimal Viable Product (MVP) features
- Internet Connection - IP Address
- How to find IP Address given Domain Name - DNS
- How to prevent ICANN from becoming a SPoF bottleneck } true like hierarchy of DNS servers
- Scaling challenges - data is too large
 - Vertical / Horizontal Scaling
 - Scale Up
 - Scale Out
- How to present a unified view to user?
 - Load Balancer

```
graph LR; DNS((DNS)) --> LB[LB]; LB --> AppServers[App Servers]; AppServers --> User
```
- Load Balancer - provide a unified view
 - distribute the load evenly across all servers
- How does LB know about what servers are available?
 - Heartbeat / Healthcheck
 - push from server to LB
 - pull from LB to servers
- How to prevent LB from becoming a SPoF & bottleneck?
 - bottleneck - LB's are able to handle higher load
 - v. large scale
 - ∴ simpler
 - SPoF → multiple LBs → DNS can store multiple IPs for a domain



API gateway / Rate Limiter / Load Balancer / Reverse Proxy - towards class end

Q1 How to route the requests? which request to which server?



We Cannot answer the routing Question until we know
↓

Q2 How is the data distributed?

① Can all the bookmark (data) be stored in 1 server?

No!

② even if we could store all data on 1 server, should we?

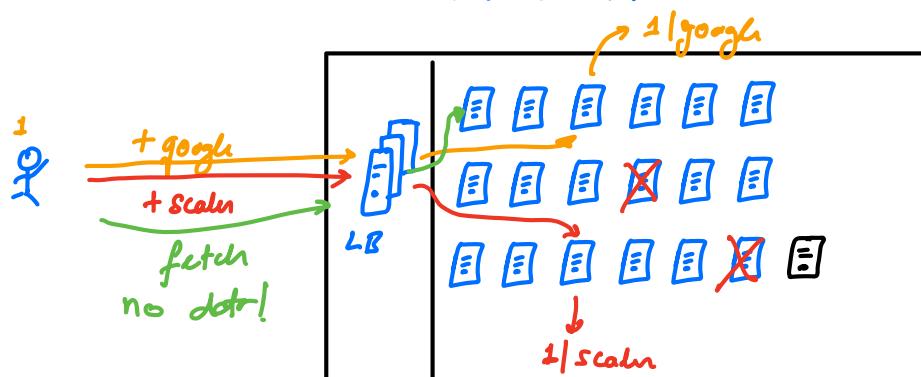
No! ↗ it will not be able to handle the load
bottleneck (# requests/sec)
It might fail spot

∴ we need to distribute the data.

③ Should I distribute data randomly?

can any (user-id / URL) go to any server? No!

No ∵ how will we retrieve it then.



∴ we should distribute data **using user-id**

All book marks of u1 go to the same server S1

" u2 " " " " " " S2

" u3 " " " " " " S3

:

Sharding

I never done randomly
always done via some sharding key
user-id (in this case study)

Users

	id	name	phone	address	username	pass-hash-salt
--	----	------	-------	---------	----------	----------------

1	Umesh
---	-------	----	----	----	----	----

2	Nidhi					
---	-------	--	--	--	--	--

3	Jai					
---	-----	--	--	--	--	--

4	Rasif					
---	-------	--	--	--	--	--

Vertical Partitioning

Normalization

(inside 1 server)

Users username pass-hash-salt

id

1

2

3

4

profiles

user-id name phone address

FKey Umes

Nidhi

Jai

Rasif

Horizontal Partitioning

Users - 1

	id	name	phone	address	username	pass-hash-salt
1	Umes
3	Jai

Users - 2

	id	name	phone	address	username	pass-hash-salt
2	Nidhi
4	Rasif

- might be within same server

↳ fast indexing

- might do this across multiple servers Sharding

Sharding - horizontally partitioning data across multiple servers

always done using a sharding-key

for this study our sharding key = user-id

Task: find a way of mapping user id → server id
 which user's data should be stored in which server?

Attempt 1

Round Robin X

N servers in total
 $= 5$

- ① is load even?
- ② fast to compute server id from user id
- X what happens if # of servers change?

U1
U6
U11

U2
U7
U12

U3
U8
U13

U4
U9
U14

U5
U10
U15

```
int N = 5; // No of servers
int getServerID(int userID) {
    return (userID % N) + 1;
}
```

$N=5$

U1
U6
U11

U2
U7
U12

U3
U8
U13

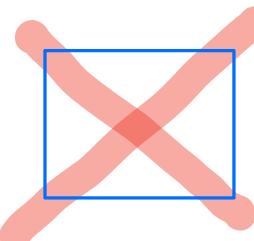
U4
U9
U14

U5
U10
U15

$N=4$

U1
U5
U9

U2
U6
U10



U3
U7
U11

U4
U8
U12

ideally only the users for whom the server has crashed should be remapped.

However here, there is a lot of unnecessary remapping.
almost all users are unmapped

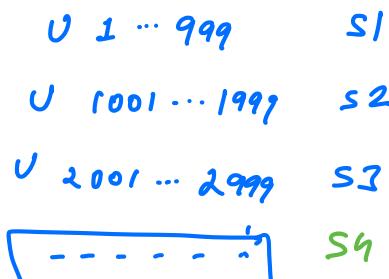
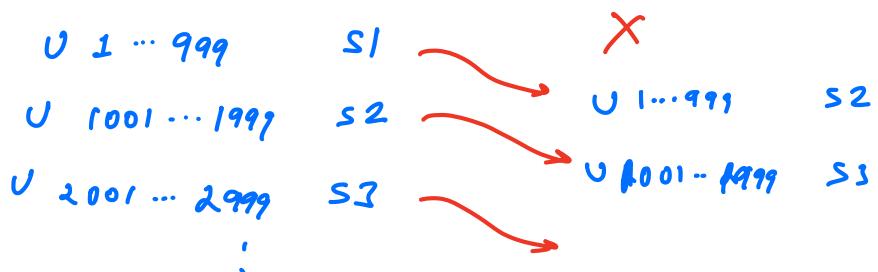
How to prevent data loss? Magic fairy (Son Pari)
backups the data using magic

similarly when we add a new server
we have lots of remapping $\therefore N$ changes

Attempt -2 Range Based X

$$\text{Server_id} = \left\lfloor \frac{\text{user_id}}{1000} \right\rfloor$$

Same problem!



Attempt 3 - Store the mapping explicitly X

Map $\langle \text{user_id}, \text{server_id} \rangle$ in RAM (initialized randomly)

if server X crashes \rightarrow find all user that were going to this server and assign new server.
No unnecessary mapping.

if a server is added \rightarrow select a % of users and move them to

this new server

issue: how to store this mapping?

① Mapping is large

60 GB

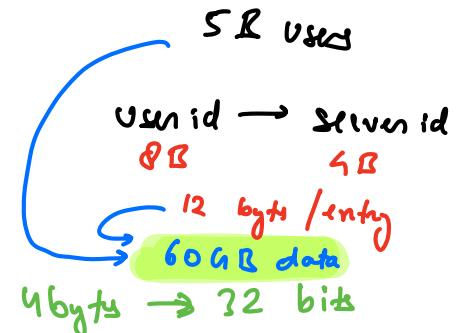
② Store in RAM?

we would like to, but

① power fails?

∴ we should store on disk

② isn't disk too slow



2^{32} possible values
≈ 4 billion

of users > 4 billion.

Youtube 4 bytes signed int
to store view count

$-2^{31} \dots 2^{31}-1$

≈ -2 billion ... 2 billion

③ what if we have multiple LBS?

How to ensure consistency (same mapping) across all LBS?

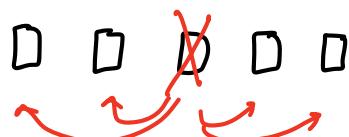
V·V·V·V·V... hard problem

impossible!!

Solution — Consistent Hashing

Goals

- ① data should be evenly distributed across servers
- ② efficiently compute user-id → server id
- ③ should not have to store any mapping data
- ④ if a server crashes each of the rest of the servers should handle part of the load



all remaining servers should share the burden of fallen server equally.

- ⑤ if a new server is added then it should relieve the load of all other servers equally



8:18 → 8:30

Hash Functions

function: mapping b/w two sets

$$f(x) = y \quad \begin{matrix} A \rightarrow B \\ \text{domain} & \text{range (codomain)} \end{matrix}$$

given x I can calculate y .

doesn't matter how many time I input x x deterministic
I will always get y y deterministic

Hash Function — seem random but deterministic

Input: anything

Output: no b/w $0 \dots K$

$$H_1(x) = e^{\sin\left(\lfloor \frac{(x^2 + 3)}{5} \rfloor\right)} \% \underbrace{(10^9 + 7)}_{(0 \dots 10^9 + 6)}$$

2713

long

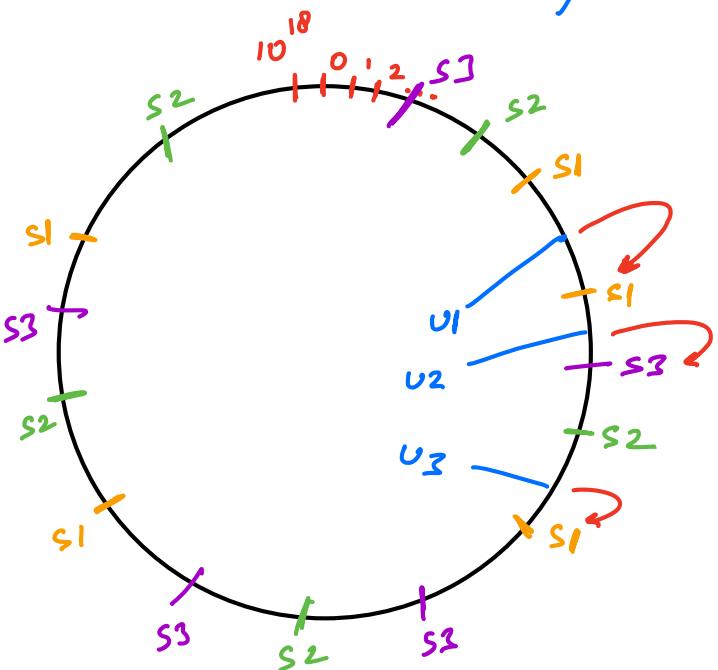
$$H_2(x) = (\text{~~~}) \% (10^9 + 7)$$

$$H_3(x) = :$$

Consistent Hashing

- ① K hash functions for each server \rightarrow virtual spot
- ② 1 hash for the user
- ③ All these hash functions have the same range

$\text{Hash}(_ _) \rightarrow$
64 bit
int



Consider $N = 3$ servers
 $S1 \quad S2 \quad S3$

$K = 5$ hashes for servers

$$\begin{aligned} H_1(S1) &= 1000 \\ H_2(S1) &= 2726 \\ H_3(S1) &= 32 \\ H_4(S1) &= 47 \\ H_5(S1) &= 110736 \end{aligned}$$

} randomly distributed
but deterministic

each user also has a spot on the ring

$$H_u(U1) = 1002$$

- ① Can the hash values collide?

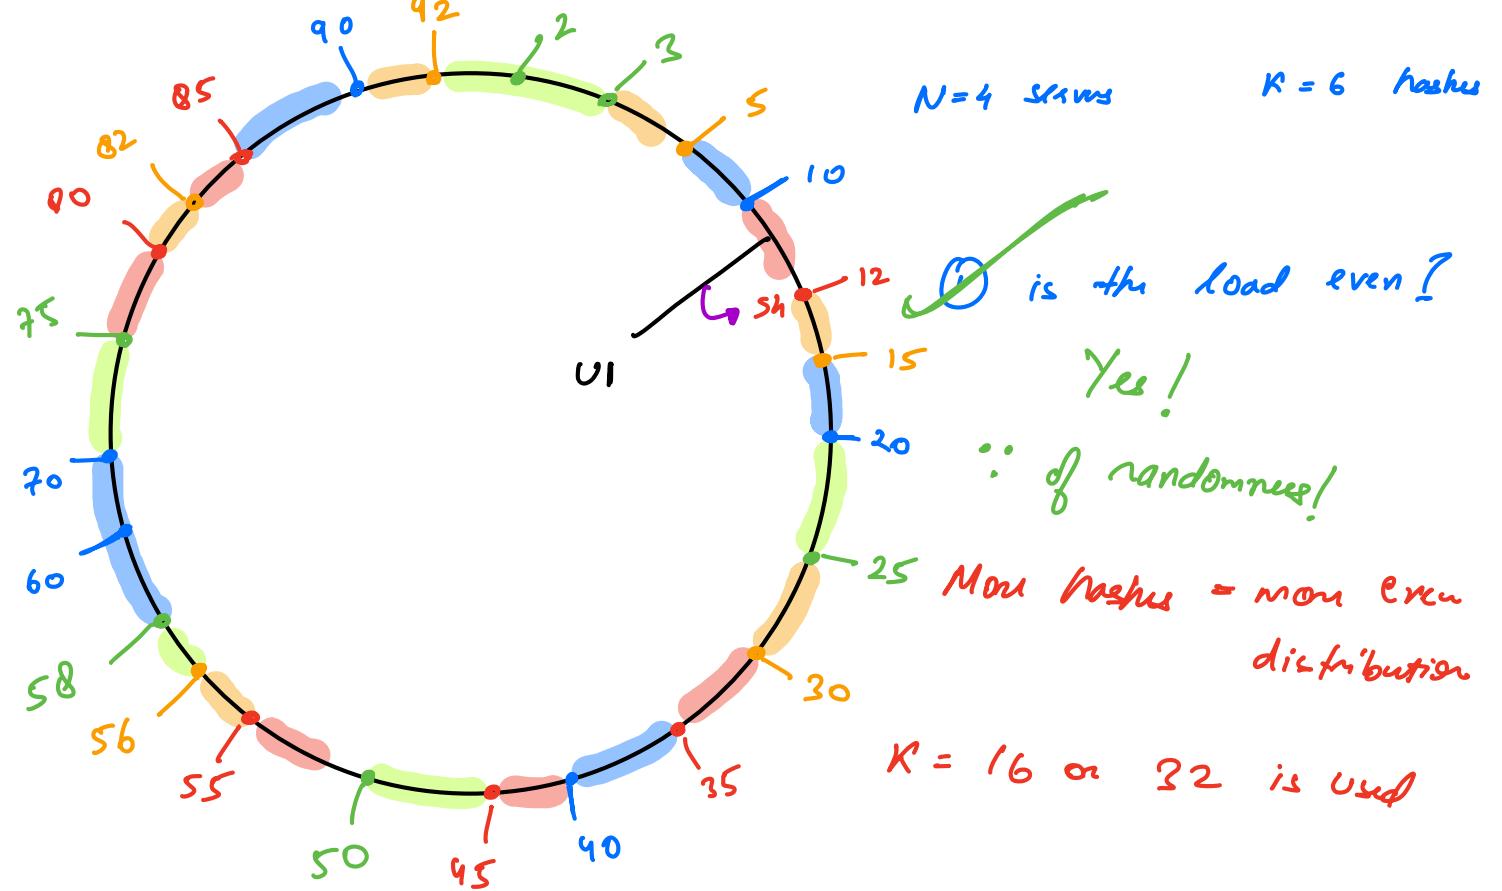
Yes! However the chances are so low that we don't have to worry!

ignore \because v. v. improbable!

& if it happens \rightarrow no issue

- ② How to route requests?

Given user-id compute the user's spot on the ring
find the first server on ring moving clockwise.



in code

$\rightarrow [$									
s_1	s_1	s_2	s_3	s_4	s_2	s_3	s_1	s_2	\dots
2	3	5	10	12	15	20	25	30	...

sorted by hash

$H_U(\text{user-id}) \rightarrow 14$ (binary search upper bound) with wrap around

① load even? Yes ∵ chance

② efficient to compute?

$$\lg_2 \left(\frac{N \cdot R}{\# q \text{ slots}} \right) = 10,000$$

③ store anything?

Yes, this sorted array

RAM? Yes ∵ < 1MB in size

$$\lg_2 (32000) \approx 15 \quad (15 \mu\text{sec / query})$$

✓ Your failure? you can always recompute this array

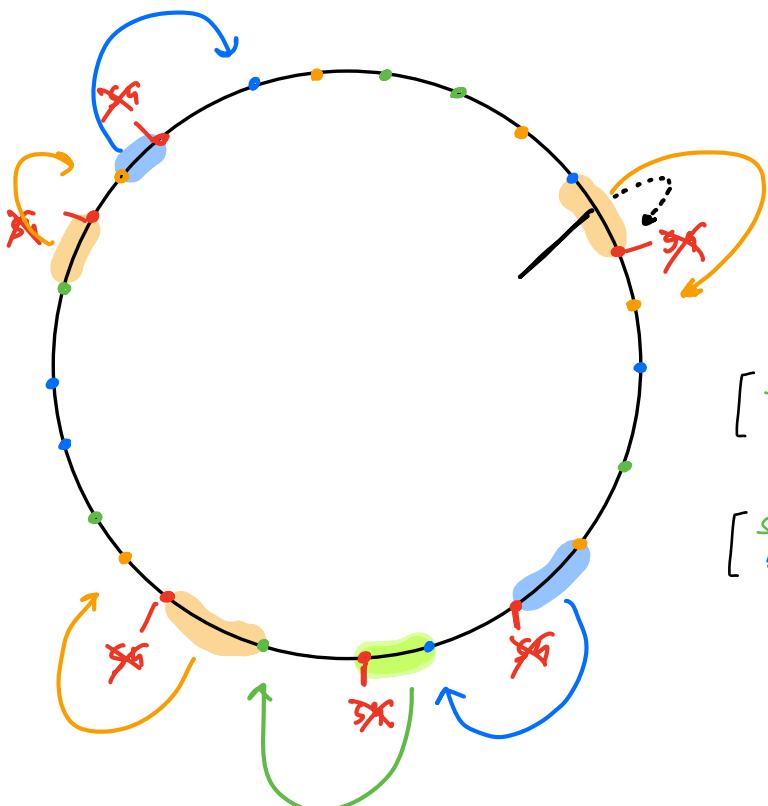
∴ hashes & servers remain same!

✓ Ensure consistency across multiple LBs?

∴ all LBs are config with the same algo

& the servers are all shared for all LBs

Server Crashes



s_4 crashes

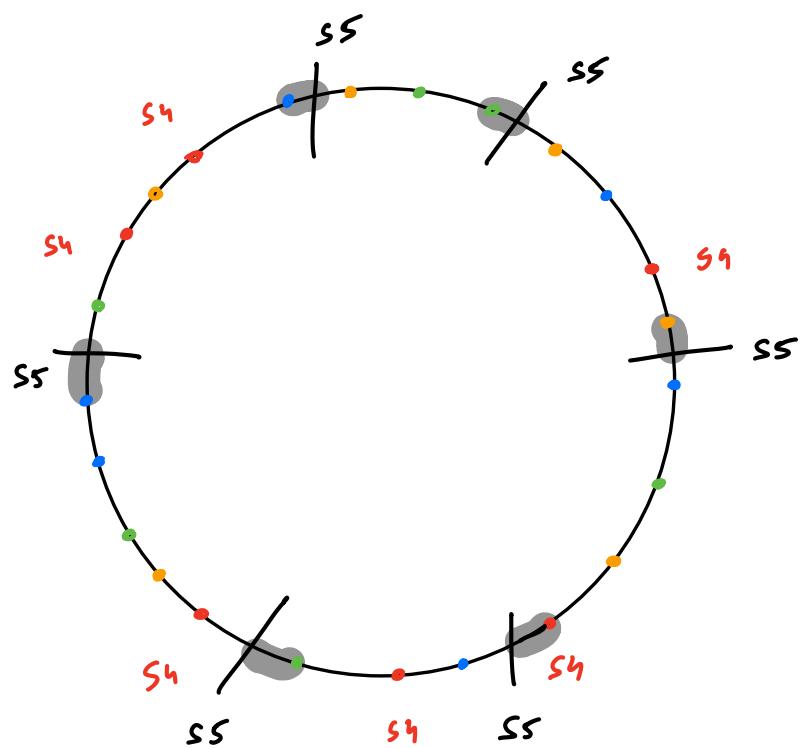
① Son Pari (Magic fairy) will somehow move data to next server

② Load of crashed server is equally handled by remaining servers

s_1	s_1	s_2	s_3	s_4	s_2	s_3	s_1	s_2	\dots
2	3	5	10	12	15	20	25	30	...

s_1	s_1	s_2	s_3	s_2	s_3	s_1	s_2	\dots
2	3	5	10	15	20	25	30	...

New Server is added



SS is added

↳ passed through it makes

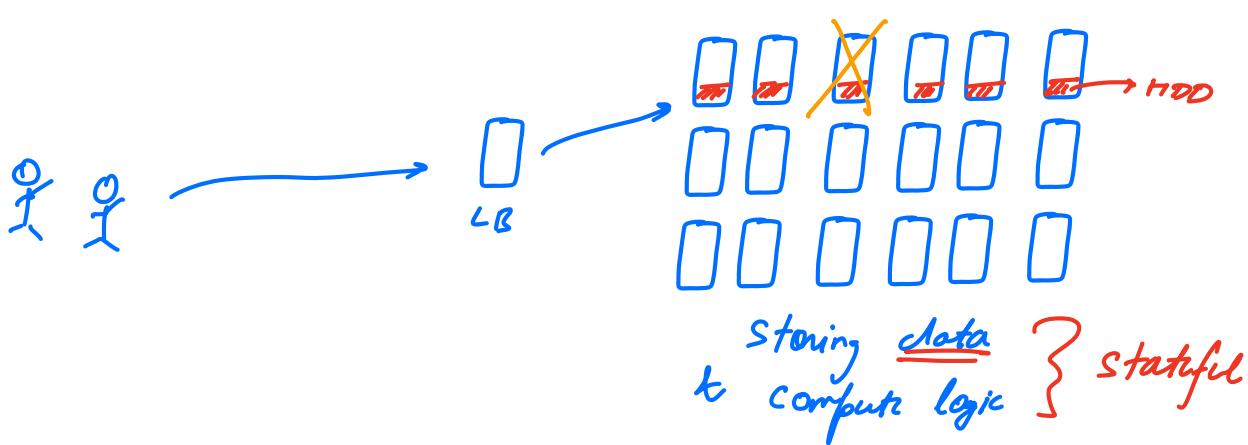
SS has taken some load
from . . .
each server!

how does SS get the data?
Son Pani *

Extremely rare that you will use consistent Hashing!

explicitly - rarely used

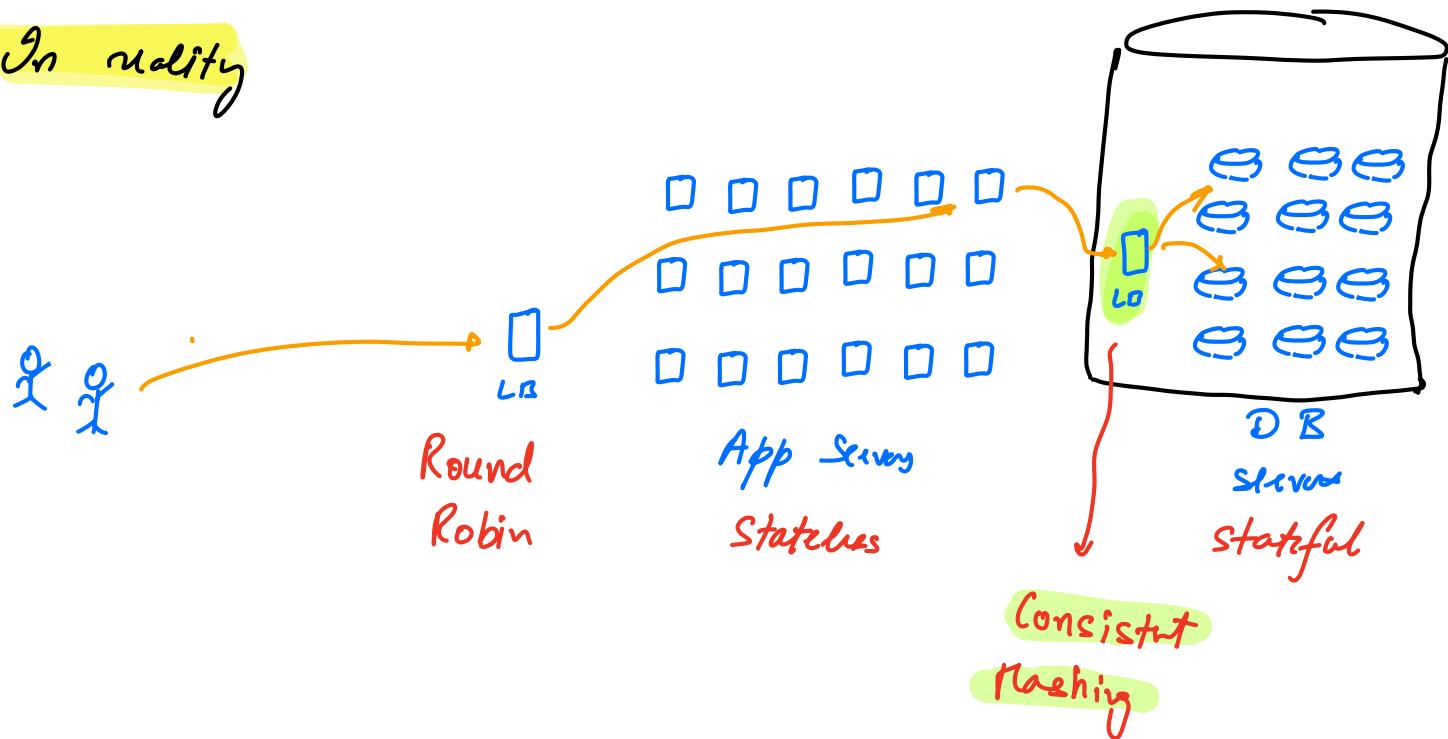
implicitly - used everywhere!!



bad architecture

∴ Coupled your data with code

In reality



State → data

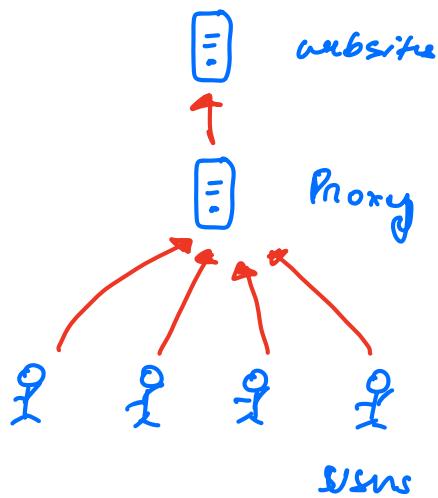
Stateful → storing data → DB is Always stateful

Stateless → no data stored

Reverse Proxy

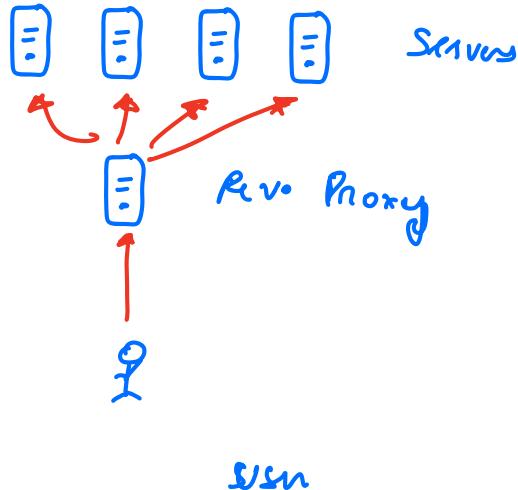
Proxy Server

hides the servers from the user

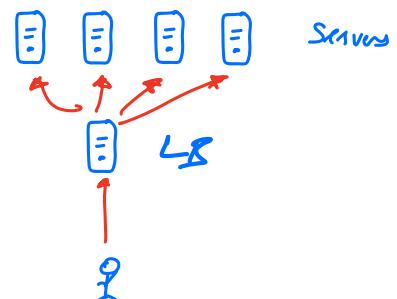


Reverse Proxy

hides the servers from the user

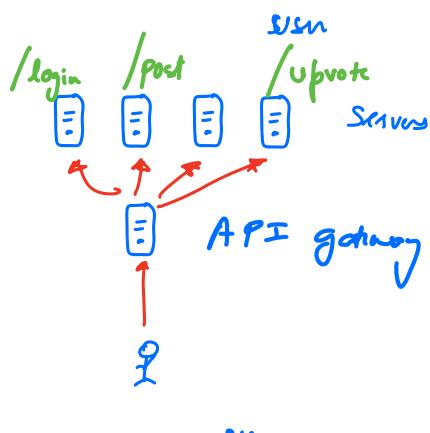


Load Balancer → unified view
distributes load evenly
is a reverse proxy

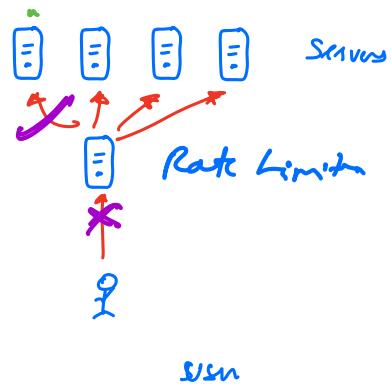


API gateway → microservices
API endpoints
routes through to the correct serv.

is also a reverse proxy



Rate Limiter → prevents abuse
↳ throttles req
is also like a reverse proxy



Nginx →
Rate Limiter
API gateway
LB

Kong / ...

store passwords in db

- Never store this in plain text

upper limit on pass length

(8-20 letters)

↓
security → only when you store
pass as plain text

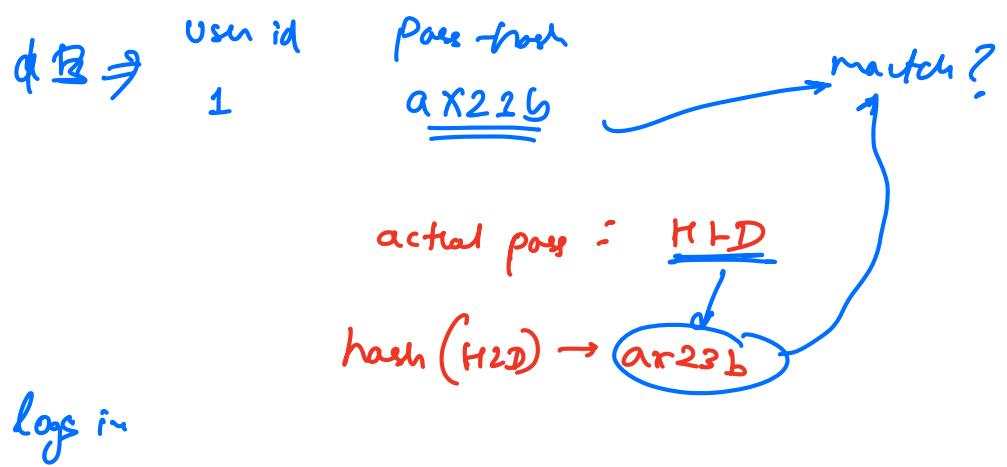
users

id	pass
	VARCHAR(20)

- hash the password

destroys the value - anyone looking at DB
cannot see the actual pass

How to validate password login?



- Anubhava & Jai both have same pass
Nanuto

User id	Pass-hash
1	ab22c
2	ab32c

= Same :: pass is same.

dictionary attack

Hash (Pass + Salt)

→ generated a random salt when user first register

→ saved in DB

dilip
user id = 10
pass = xyz
salt = 12ab3#23
(random)

hash (xyz 12ab3#23)

db user login

bcrypt

Encryption

\neq Hash

inversible

reversible

lossless

lossy

Never find original data